


SOLVED PRACTICE WORKBOOK

Artificial Intelligence Governance and IndiaAI - Solved Practice Workbook

UPSC complete learning session | Foundation to advanced | Concepts, evidence, practice and answer writing



Original deterministic study visual; labels are explanatory, not textual quotations.

Source order: repository knowledge -> official current linkage -> clearly marked analysis. No fabricated PYQs or statistics.

CONTENTS / WORKBOOK INDEX

Major learning parts and meaningful subtopics are listed in document order. Page numbers are generated from the final PDF layout; PDF bookmarks mirror the same hierarchy.

Artificial Intelligence Governance and IndiaAI - Solved Practice Workbook 5

BASIC MCQS / REMEDIATION 5

Q1. Which statement correctly identifies AI-taxonomy boundary?	5
Q2. Which option preserves the technical boundary of AI-taxonomy boundary?	5
Q3. Which statement uses AI-taxonomy boundary without changing its institution, unit or status?	5
Q4. Which option avoids the standard UPSC close-option trap about AI-taxonomy boundary?	6
Q5. Which statement correctly identifies Capability-deployment boundary?	6
Q6. Which option preserves the technical boundary of Capability-deployment boundary?	6
Q7. Which statement uses Capability-deployment boundary without changing its institution, unit or status?	6
Q8. Which option avoids the standard UPSC close-option trap about Capability-deployment boundary?	7
Q9. Which statement correctly identifies Prediction-causation boundary?	7
Q10. Which option preserves the technical boundary of Prediction-causation boundary?	7
Q11. Which statement uses Prediction-causation boundary without changing its institution, unit or status?	8
Q12. Which option avoids the standard UPSC close-option trap about Prediction-causation boundary?	8
Q13. Which statement correctly identifies Foundation-model boundary?	8
Q14. Which option preserves the technical boundary of Foundation-model boundary?	8
Q15. Which statement uses Foundation-model boundary without changing its institution, unit or status?	9
Q16. Which option avoids the standard UPSC close-option trap about Foundation-model boundary?	9
Q17. Which statement correctly identifies Hallucination-boundary?	9
Q18. Which option preserves the technical boundary of Hallucination-boundary?	10
Q19. Which statement uses Hallucination-boundary without changing its institution, unit or status?	10
Q20. Which option avoids the standard UPSC close-option trap about Hallucination-boundary?	10
Q21. Which statement correctly identifies IndiaAI-mission boundary?	10
Q22. Which option preserves the technical boundary of IndiaAI-mission boundary?	11
Q23. Which statement uses IndiaAI-mission boundary without changing its institution, unit or status?	11
Q24. Which option avoids the standard UPSC close-option trap about IndiaAI-mission boundary?	11
Q25. Which statement correctly identifies Implementation-institution boundary?	12
Q26. Which option preserves the technical boundary of Implementation-institution boundary?	12
Q27. Which statement uses Implementation-institution boundary without changing its institution, unit or status?	12

Q28. Which option avoids the standard UPSC close-option trap about Implementation-institution boundary?	12
Q29. Which statement correctly identifies Pillar-status boundary?	13
Q30. Which option preserves the technical boundary of Pillar-status boundary?	13
Q31. Which statement uses Pillar-status boundary without changing its institution, unit or status?	13
Q32. Which option avoids the standard UPSC close-option trap about Pillar-status boundary?	14
Q33. Which statement correctly identifies Compute-measure boundary?	14
Q34. Which option preserves the technical boundary of Compute-measure boundary?	14
Q35. Which statement uses Compute-measure boundary without changing its institution, unit or status?	14
Q36. Which option avoids the standard UPSC close-option trap about Compute-measure boundary?	15
Q37. Which statement correctly identifies AIKosh-data boundary?	15
Q38. Which option preserves the technical boundary of AIKosh-data boundary?	15
Q39. Which statement uses AIKosh-data boundary without changing its institution, unit or status?	16
Q40. Which option avoids the standard UPSC close-option trap about AIKosh-data boundary?	16
Q41. Which statement correctly identifies Use-case-model boundary?	16
Q42. Which option preserves the technical boundary of Use-case-model boundary?	16
Q43. Which statement uses Use-case-model boundary without changing its institution, unit or status?	17
Q44. Which option avoids the standard UPSC close-option trap about Use-case-model boundary?	17
Q45. Which statement correctly identifies Decision-support boundary?	17
Q46. Which option preserves the technical boundary of Decision-support boundary?	18
Q47. Which statement uses Decision-support boundary without changing its institution, unit or status?	18
Q48. Which option avoids the standard UPSC close-option trap about Decision-support boundary?	18
Q49. Which statement correctly identifies Risk-lifecycle boundary?	19
Q50. Which option preserves the technical boundary of Risk-lifecycle boundary?	19
Q51. Which statement uses Risk-lifecycle boundary without changing its institution, unit or status?	19
Q52. Which option avoids the standard UPSC close-option trap about Risk-lifecycle boundary?	19
Q53. Which statement correctly identifies Bias-fairness boundary?	20
Q54. Which option preserves the technical boundary of Bias-fairness boundary?	20
Q55. Which statement uses Bias-fairness boundary without changing its institution, unit or status?	20
Q56. Which option avoids the standard UPSC close-option trap about Bias-fairness boundary?	21
Q57. Which statement correctly identifies Privacy-security boundary?	21
Q58. Which option preserves the technical boundary of Privacy-security boundary?	21
Q59. Which statement uses Privacy-security boundary without changing its institution, unit or status?	21

Q60. Which option avoids the standard UPSC close-option trap about Privacy-security boundary?	22
Q61. Which statement correctly identifies Deepfake-instrument boundary?	22
Q62. Which option preserves the technical boundary of Deepfake-instrument boundary?	22
Q63. Which statement uses Deepfake-instrument boundary without changing its institution, unit or status?	23
Q64. Which option avoids the standard UPSC close-option trap about Deepfake-instrument boundary?	23
Q65. Which statement correctly identifies Guideline-law boundary?	23
Q66. Which option preserves the technical boundary of Guideline-law boundary?	24
Q67. Which statement uses Guideline-law boundary without changing its institution, unit or status?	24
Q68. Which option avoids the standard UPSC close-option trap about Guideline-law boundary?	24
Q69. Which statement correctly identifies Procurement-accountability boundary?	24
Q70. Which option preserves the technical boundary of Procurement-accountability boundary?	25
Q71. Which statement uses Procurement-accountability boundary without changing its institution, unit or status?	25
Q72. Which option avoids the standard UPSC close-option trap about Procurement-accountability boundary?	25
Q73. Which statement correctly identifies Sectoral-governance boundary?	26
Q74. Which option preserves the technical boundary of Sectoral-governance boundary?	26
Q75. Which statement uses Sectoral-governance boundary without changing its institution, unit or status?	26
Q76. Which option avoids the standard UPSC close-option trap about Sectoral-governance boundary?	26
Q77. Which statement correctly identifies Volatile-mission boundary?	27
Q78. Which option preserves the technical boundary of Volatile-mission boundary?	27
Q79. Which statement uses Volatile-mission boundary without changing its institution, unit or status?	27
Q80. Which option avoids the standard UPSC close-option trap about Volatile-mission boundary?	28

PYQS AND ANSWER PRACTICE 28

AUDITED TOPIC-SPECIFIC PYQ OWNERSHIP	28
PYQ DEMAND CARD 1 - 2020 Prelims GS-I	28
PYQ DEMAND CARD 2 - 2023 GS-III	29
PYQ DEMAND CARD 3 - 2025 GS-I	30
ORIGINAL MAINS 1 - 10 MARKS	32
ORIGINAL MAINS 2 - 10 MARKS	34
ORIGINAL MAINS 3 - 15 MARKS	36
ORIGINAL MAINS 4 - 15 MARKS	40
ORIGINAL MAINS 5 - 20 MARKS	43
ORIGINAL MAINS 6 - 20 MARKS	46

Artificial Intelligence Governance and IndiaAI - Solved Practice Workbook

BASIC MCQS / REMEDIATION

Q1. Which statement correctly identifies AI-taxonomy boundary?

A. Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. B. A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. C. AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. D. A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it.

Answer: A. Explanation: Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q2. Which option preserves the technical boundary of AI-taxonomy boundary?

A. Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. B. Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. C. AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. D. A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it.

Answer: B. Explanation: Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q3. Which statement uses AI-taxonomy boundary without changing its institution, unit or status?

A. IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. B. Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. C. Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. D. A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it.

Answer: C. Explanation: Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q4. Which option avoids the standard UPSC close-option trap about AI-taxonomy boundary?

A. Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. B. IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. C. MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. D. Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task.

Answer: D. Explanation: Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q5. Which statement correctly identifies Capability-deployment boundary?

A. A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. B. AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. C. Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. D. A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it.

Answer: A. Explanation: A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q6. Which option preserves the technical boundary of Capability-deployment boundary?

A. Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. B. A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. C. A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. D. IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code.

Answer: B. Explanation: A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q7. Which statement uses Capability-deployment boundary without changing its institution, unit or status?

A. Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. B. IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. C. A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. D. MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making.

Answer: C. Explanation: A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q8. Which option avoids the standard UPSC close-option trap about Capability-deployment boundary?

A. A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. B. IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. C. MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. D. A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting.

Answer: D. Explanation: A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q9. Which statement correctly identifies Prediction-causation boundary?

A. AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. B. Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. C. IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. D. A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it.

Answer: A. Explanation: AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q10. Which option preserves the technical boundary of Prediction-causation boundary?

A. Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. B. AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. C. IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. D. MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making.

Answer: B. Explanation: AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q11. Which statement uses Prediction-causation boundary without changing its institution, unit or status?

A. A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. B. IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. C. AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. D. MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making.

Answer: C. Explanation: AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q12. Which option avoids the standard UPSC close-option trap about Prediction-causation boundary?

A. MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. B. A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. C. A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. D. AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification.

Answer: D. Explanation: AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q13. Which statement correctly identifies Foundation-model boundary?

A. A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. B. MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. C. IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. D. Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary.

Answer: A. Explanation: A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q14. Which option preserves the technical boundary of Foundation-model boundary?

A. MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. B. A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. C. A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. D. IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code.

Answer: B. Explanation: A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q15. Which statement uses Foundation-model boundary without changing its institution, unit or status?

A. A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. B. A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. C. A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. D. MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making.

Answer: C. Explanation: A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q16. Which option avoids the standard UPSC close-option trap about Foundation-model boundary?

A. A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. B. AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. C. A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. D. A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it.

Answer: D. Explanation: A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q17. Which statement correctly identifies Hallucination-boundary?

A. Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. B. MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. C. A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. D. IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code.

Answer: A. Explanation: Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q18. Which option preserves the technical boundary of Hallucination-boundary?

A. A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. B. Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. C. MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. D. A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition.

Answer: B. Explanation: Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q19. Which statement uses Hallucination-boundary without changing its institution, unit or status?

A. A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. B. A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. C. Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. D. AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access.

Answer: C. Explanation: Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q20. Which option avoids the standard UPSC close-option trap about Hallucination-boundary?

A. AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. B. A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. C. An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. D. Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary.

Answer: D. Explanation: Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q21. Which statement correctly identifies IndiaAI-mission boundary?

A. IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. B. MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. C. A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. D. A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence.

Answer: A. Explanation: IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q22. Which option preserves the technical boundary of IndiaAI-mission boundary?

A. AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. B. IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. C. A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. D. A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition.

Answer: B. Explanation: IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q23. Which statement uses IndiaAI-mission boundary without changing its institution, unit or status?

A. A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. B. An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. C. IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. D. AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access.

Answer: C. Explanation: IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q24. Which option avoids the standard UPSC close-option trap about IndiaAI-mission boundary?

A. AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. B. An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. C. AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. D. IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code.

Answer: D. Explanation: IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q25. Which statement correctly identifies Implementation-institution boundary?

A. MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. B. A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. C. A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. D. AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access.

Answer: A. Explanation: MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q26. Which option preserves the technical boundary of Implementation-institution boundary?

A. AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. B. MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. C. A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. D. An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress.

Answer: B. Explanation: MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q27. Which statement uses Implementation-institution boundary without changing its institution, unit or status?

A. An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. B. AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. C. MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. D. AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access.

Answer: C. Explanation: MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q28. Which option avoids the standard UPSC close-option trap about Implementation-institution boundary?

A. AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. B. An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. C. AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. D. MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making.

Answer: D. Explanation: MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q29. Which statement correctly identifies Pillar-status boundary?

A. A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. B. A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. C. An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. D. AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access.

Answer: A. Explanation: A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q30. Which option preserves the technical boundary of Pillar-status boundary?

A. AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. B. A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. C. An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. D. AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal.

Answer: B. Explanation: A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q31. Which statement uses Pillar-status boundary without changing its institution, unit or status?

A. AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. B. An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. C. A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. D. AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal.

Answer: C. Explanation: A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q32. Which option avoids the standard UPSC close-option trap about Pillar-status boundary?

A. Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. B. AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. C. AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. D. A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence.

Answer: D. Explanation: A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q33. Which statement correctly identifies Compute-measure boundary?

A. A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. B. An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. C. AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. D. AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal.

Answer: A. Explanation: A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q34. Which option preserves the technical boundary of Compute-measure boundary?

A. AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. B. A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. C. An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. D. AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal.

Answer: B. Explanation: A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q35. Which statement uses Compute-measure boundary without changing its institution, unit or status?

A. Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. B. AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. C. A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. D. AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal.

Answer: C. Explanation: A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q36. Which option avoids the standard UPSC close-option trap about Compute-measure boundary?

A. Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. B. AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. C. Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. D. A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition.

Answer: D. Explanation: A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q37. Which statement correctly identifies AIKosh-data boundary?

A. AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. B. An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. C. AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. D. AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal.

Answer: A. Explanation: AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q38. Which option preserves the technical boundary of AIKosh-data boundary?

A. AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. B. AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. C. Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. D. AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle.

Answer: B. Explanation: AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q39. Which statement uses AIKosh-data boundary without changing its institution, unit or status?

A. Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. B. AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. C. AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. D. Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability.

Answer: C. Explanation: AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q40. Which option avoids the standard UPSC close-option trap about AIKosh-data boundary?

A. Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. B. Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. C. Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. D. AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access.

Answer: D. Explanation: AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q41. Which statement correctly identifies Use-case-model boundary?

A. An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. B. AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. C. Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. D. AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal.

Answer: A. Explanation: An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q42. Which option preserves the technical boundary of Use-case-model boundary?

A. Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. B. An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. C. AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. D. Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability.

Answer: B. Explanation: An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q43. Which statement uses Use-case-model boundary without changing its institution, unit or status?

A. Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. B. Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. C. An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. D. Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime.

Answer: C. Explanation: An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q44. Which option avoids the standard UPSC close-option trap about Use-case-model boundary?

A. An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. B. Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. C. Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. D. An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress.

Answer: D. Explanation: An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q45. Which statement correctly identifies Decision-support boundary?

A. AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. B. Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. C. Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. D. AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle.

Answer: A. Explanation: AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q46. Which option preserves the technical boundary of Decision-support boundary?

A. Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. B. AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. C. Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. D. Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy.

Answer: B. Explanation: AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q47. Which statement uses Decision-support boundary without changing its institution, unit or status?

A. An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. B. Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. C. AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. D. Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability.

Answer: C. Explanation: AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q48. Which option avoids the standard UPSC close-option trap about Decision-support boundary?

A. Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. B. Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. C. An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. D. AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal.

Answer: D. Explanation: AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q49. Which statement correctly identifies Risk-lifecycle boundary?

A. AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. B. Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. C. Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. D. Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability.

Answer: A. Explanation: AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q50. Which option preserves the technical boundary of Risk-lifecycle boundary?

A. Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. B. AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. C. An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. D. Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime.

Answer: B. Explanation: AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q51. Which statement uses Risk-lifecycle boundary without changing its institution, unit or status?

A. An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. B. Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. C. AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. D. Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime.

Answer: C. Explanation: AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q52. Which option avoids the standard UPSC close-option trap about Risk-lifecycle boundary?

A. An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. B. AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. C. Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. D. AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle.

Answer: D. Explanation: AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q53. Which statement correctly identifies Bias-fairness boundary?

A. Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. B. An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. C. Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. D. Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime.

Answer: A. Explanation: Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q54. Which option preserves the technical boundary of Bias-fairness boundary?

A. Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. B. Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. C. Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. D. An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record.

Answer: B. Explanation: Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q55. Which statement uses Bias-fairness boundary without changing its institution, unit or status?

A. AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. B. An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. C. Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. D. Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability.

Answer: C. Explanation: Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q56. Which option avoids the standard UPSC close-option trap about Bias-fairness boundary?

A. AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. B. Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. C. Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. D. Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy.

Answer: D. Explanation: Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q57. Which statement correctly identifies Privacy-security boundary?

A. Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. B. Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. C. An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. D. Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime.

Answer: A. Explanation: Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q58. Which option preserves the technical boundary of Privacy-security boundary?

A. An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. B. Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. C. AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. D. Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability.

Answer: B. Explanation: Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q59. Which statement uses Privacy-security boundary without changing its institution, unit or status?

A. Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. B. AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. C. Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. D. Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability.

Answer: C. Explanation: Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q60. Which option avoids the standard UPSC close-option trap about Privacy-security boundary?

A. Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. B. Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. C. AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. D. Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability.

Answer: D. Explanation: Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q61. Which statement correctly identifies Deepfake-instrument boundary?

A. Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. B. AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. C. An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. D. Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability.

Answer: A. Explanation: Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q62. Which option preserves the technical boundary of Deepfake-instrument boundary?

A. Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. B. Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. C. AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. D. Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability.

Answer: B. Explanation: Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q63. Which statement uses Deepfake-instrument boundary without changing its institution, unit or status?

A. AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. B. Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. C. Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. D. Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability.

Answer: C. Explanation: Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q64. Which option avoids the standard UPSC close-option trap about Deepfake-instrument boundary?

A. A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. B. Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. C. Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. D. Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime.

Answer: D. Explanation: Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q65. Which statement correctly identifies Guideline-law boundary?

A. An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. B. AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. C. Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. D. Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability.

Answer: A. Explanation: An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q66. Which option preserves the technical boundary of Guideline-law boundary?

A. Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. B. An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. C. Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. D. AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators.

Answer: B. Explanation: An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q67. Which statement uses Guideline-law boundary without changing its institution, unit or status?

A. Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. B. Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. C. An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. D. A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting.

Answer: C. Explanation: An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q68. Which option avoids the standard UPSC close-option trap about Guideline-law boundary?

A. Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. B. A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. C. AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. D. An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record.

Answer: D. Explanation: An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q69. Which statement correctly identifies Procurement-accountability boundary?

A. Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. B. Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. C. AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. D. Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability.

Answer: A. Explanation: Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q70. Which option preserves the technical boundary of Procurement-accountability boundary?

A. Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. B. Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. C. Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. D. A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting.

Answer: B. Explanation: Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q71. Which statement uses Procurement-accountability boundary without changing its institution, unit or status?

A. A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. B. AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. C. Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. D. Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task.

Answer: C. Explanation: Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q72. Which option avoids the standard UPSC close-option trap about Procurement-accountability boundary?

A. AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. B. A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. C. A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. D. Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability.

Answer: D. Explanation: Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q73. Which statement correctly identifies Sectoral-governance boundary?

A. AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. B. Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. C. A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. D. Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability.

Answer: A. Explanation: AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q74. Which option preserves the technical boundary of Sectoral-governance boundary?

A. A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. B. AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. C. Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. D. AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification.

Answer: B. Explanation: AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q75. Which statement uses Sectoral-governance boundary without changing its institution, unit or status?

A. AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. B. A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. C. AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. D. A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting.

Answer: C. Explanation: AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q76. Which option avoids the standard UPSC close-option trap about Sectoral-governance boundary?

A. AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. B. Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. C. A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. D. AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators.

Answer: D. Explanation: AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q77. Which statement correctly identifies Volatile-mission boundary?

A. Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. B. AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. C. Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. D. A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting.

Answer: A. Explanation: Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q78. Which option preserves the technical boundary of Volatile-mission boundary?

A. A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. B. Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. C. AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. D. A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting.

Answer: B. Explanation: Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q79. Which statement uses Volatile-mission boundary without changing its institution, unit or status?

A. Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. B. A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. C. Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. D. AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification.

Answer: C. Explanation: Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

Q80. Which option avoids the standard UPSC close-option trap about Volatile-mission boundary?

A. IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. B. A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. C. Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. D. Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability.

Answer: D. Explanation: Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. The other options change the scientific category, institutional owner, measurement, mission rung or operating status.

PYQS AND ANSWER PRACTICE

AUDITED TOPIC-SPECIFIC PYQ OWNERSHIP

Audited ledgers route the 2020 AI-capability objective concept, the 2023 clinical-diagnosis/privacy Mains demand and the 2025 AI-drones-GIS planning demand here. No objective key, model claim or summit outcome is invented.

PYQ DEMAND CARD 1 - 2020 Prelims GS-I

Demand: Assess current AI capabilities across industry and society.

Status: Verified routed objective concept; the official key was unavailable locally, so no answer letter is asserted.

Model solution: AI-taxonomy boundary: Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. **Capability-deployment boundary:** A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Prediction-causation boundary:** AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. **Use-case-model boundary:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. The solution therefore answers the routed demand without inventing an official model answer or an objective answer key.

Demand decoding: Treat "PYQ DEMAND CARD 1 - 2020 Prelims GS-I" as a mechanism, category, institution, unit, date and status problem.

Detailed examiner-grade model answer:

Introduction and thesis: AI-taxonomy boundary: Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. **Capability-deployment boundary:** A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Prediction-causation boundary:** AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. **Use-case-model boundary:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. The solution therefore answers the routed demand without inventing an official model answer or an objective answer key.

Analytical body:

1. **Claim and named evidence:** Demand: Assess current AI capabilities across industry and society. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty,

exception or residual risk.

2. **Claim and named evidence:** Status: Verified routed objective concept; the official key was unavailable locally, so no answer letter is asserted. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.

Counter-position / limit: A launch, test, target, budget, capacity, approval, contract, dataset, benchmark or prototype cannot alone establish deployment, induction, commercial operation, safety, inclusion or observed outcome; test mechanism, status, scale and evidence.

Qualified conclusion: AI-taxonomy boundary: Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. **Capability-deployment boundary:** A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Prediction-causation boundary:** AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. **Use-case-model boundary:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. The solution therefore answers the routed demand without inventing an official model answer or an objective answer key.

Executable exam-length answer / compression plan: Identify the object and actor; trace the technical mechanism; fix the measurement and platform; test each statement against source date, status rung and closest exception.

Why this earns marks: It prevents organisation, platform, process, application, regulatory role, target and observed outcome from being conflated.

How to improve this answer: For “PYQ DEMAND CARD 1 - 2020 Prelims GS-I”, explain why the nearest distractor fails on mechanism, platform, unit, institution, legal status, maturity or date.

PYQ DEMAND CARD 2 - 2023 GS-III

Demand: Discuss AI in clinical diagnosis and threats to individual privacy.

Status: Verified routed Mains demand; the solution separates decision support, privacy, safety and clinician accountability.

Model solution: Capability-deployment boundary: A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Use-case-model boundary:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Decision-support boundary:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Privacy-security boundary:** Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. **Sectoral-governance boundary:** AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. The solution therefore answers the routed demand without inventing an official model answer or an objective answer key.

Demand decoding: The directive **answer** requires a direct position on “PYQ DEMAND CARD 2 - 2023 GS-III”, each clause, the exact technical mechanism, named Indian institution, dated status, application, risk and qualification.

Detailed examiner-grade model answer:

Introduction and thesis: Capability-deployment boundary: A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Use-case-model boundary:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Decision-support boundary:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons,

human oversight, contestability and appeal. **Privacy-security boundary:** Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. **Sectoral-governance boundary:** AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. The solution therefore answers the routed demand without inventing an official model answer or an objective answer key.

Analytical body:

1. **Claim and named evidence:** Demand: Discuss AI in clinical diagnosis and threats to individual privacy. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
2. **Claim and named evidence:** Status: Verified routed Mains demand; the solution separates decision support, privacy, safety and clinician accountability. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.

Counter-position / limit: A launch, test, target, budget, capacity, approval, contract, dataset, benchmark or prototype cannot alone establish deployment, induction, commercial operation, safety, inclusion or observed outcome; test mechanism, status, scale and evidence.

Qualified conclusion: Capability-deployment boundary: A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Use-case-model boundary:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Decision-support boundary:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Privacy-security boundary:** Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. **Sectoral-governance boundary:** AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. The solution therefore answers the routed demand without inventing an official model answer or an objective answer key.

Executable exam-length answer / compression plan: For 15 marks, spend one-sixth of the time decoding the directive and drawing definition -> mechanism -> system/institution -> application/status -> risk/outcome; write four to seven claim -> named evidence -> analysis -> qualification points; reserve the final minute for unit, date, actor, platform, approval/test/deployment rung and residual-risk checks.

Why this earns marks: The answer obeys the directive, explains science rather than listing schemes, uses India-centric evidence and preserves technical, institutional, regulatory and maturity distinctions.

How to improve this answer: For "PYQ DEMAND CARD 2 - 2023 GS-III", replace the weakest catalogue point with one exact mechanism, named mission/institution/platform, dated status, measurable boundary, implementation constraint and answer-specific qualification.

PYQ DEMAND CARD 3 - 2025 GS-I

Demand: Discuss AI and drones with GIS and remote sensing in locational and areal planning.

Status: Verified routed Mains demand; the solution treats AI as decision support and retains validation, ground-truth and public-accountability limits.

Model solution: Capability-deployment boundary: A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Prediction-causation boundary:** AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. **Use-case-model boundary:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance

redress. **Decision-support boundary:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Risk-lifecycle boundary:** AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. The solution therefore answers the routed demand without inventing an official model answer or an objective answer key.

Demand decoding: The directive **answer** requires a direct position on “PYQ DEMAND CARD 3 - 2025 GS-I”, each clause, the exact technical mechanism, named Indian institution, dated status, application, risk and qualification.

Detailed examiner-grade model answer:

Introduction and thesis: Capability-deployment boundary: A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Prediction-causation boundary:** AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. **Use-case-model boundary:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Decision-support boundary:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Risk-lifecycle boundary:** AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. The solution therefore answers the routed demand without inventing an official model answer or an objective answer key.

Analytical body:

1. **Claim and named evidence:** Demand: Discuss AI and drones with GIS and remote sensing in locational and areal planning. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
2. **Claim and named evidence:** Status: Verified routed Mains demand; the solution treats AI as decision support and retains validation, ground-truth and public-accountability limits. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.

Counter-position / limit: A launch, test, target, budget, capacity, approval, contract, dataset, benchmark or prototype cannot alone establish deployment, induction, commercial operation, safety, inclusion or observed outcome; test mechanism, status, scale and evidence.

Qualified conclusion: Capability-deployment boundary: A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Prediction-causation boundary:** AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. **Use-case-model boundary:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Decision-support boundary:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Risk-lifecycle boundary:** AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. The solution therefore answers the routed demand without inventing an official model answer or an objective answer key.

Executable exam-length answer / compression plan: For 15 marks, spend one-sixth of the time decoding the directive and drawing definition -> mechanism -> system/institution -> application/status -> risk/outcome; write four to seven claim -> named evidence -> analysis -> qualification points; reserve the final minute for unit, date, actor, platform, approval/test/deployment rung and residual-risk checks.

Why this earns marks: The answer obeys the directive, explains science rather than listing schemes, uses India-centric evidence and preserves technical, institutional, regulatory and maturity distinctions.

How to improve this answer: For “PYQ DEMAND CARD 3 - 2025 GS-I”, replace the weakest catalogue point with one exact mechanism, named mission/institution/platform, dated status, measurable boundary, implementation constraint and answer-specific qualification.

ORIGINAL MAINS 1 - 10 MARKS

Question: Distinguish AI, machine learning, deep learning and generative AI. Answer in about 150 words.

Model thesis: **Claim:** AI-taxonomy boundary. **Named evidence/example:** Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication.

Qualification: The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Prediction-causation boundary. **Named evidence/example:** AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Foundation-model boundary. **Named evidence/example:** A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Hallucination-boundary. **Named evidence/example:** Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Claim -> named evidence -> analysis -> qualification:

- Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task.
- AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification.
- A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it.
- Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary.

Qualified conclusion: **Claim:** AI-taxonomy boundary. **Named evidence/example:** Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication.

Qualification: The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Prediction-causation boundary. **Named evidence/example:** AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Foundation-model boundary. **Named evidence/example:** A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion

retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Hallucination-boundary. **Named evidence/example:** Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Demand decoding: The directive **answer** requires a direct position on "Distinguish AI, machine learning, deep learning and generative AI. Answer in about 150 words.", each clause, the exact technical mechanism, named Indian institution, dated status, application, risk and qualification.

Detailed examiner-grade model answer:

Introduction and thesis: **Claim:** AI-taxonomy boundary. **Named evidence/example:** Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Prediction-causation boundary. **Named evidence/example:** AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Foundation-model boundary. **Named evidence/example:** A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Hallucination-boundary. **Named evidence/example:** Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Analytical body:

1. **Claim and named evidence:** Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
2. **Claim and named evidence:** AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
3. **Claim and named evidence:** A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
4. **Claim and named evidence:** Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.

Counter-position / limit: A launch, test, target, budget, capacity, approval, contract, dataset, benchmark or prototype cannot alone establish deployment, induction, commercial operation, safety, inclusion or observed outcome; test mechanism, status, scale and evidence.

Qualified conclusion: Claim: AI-taxonomy boundary. **Named evidence/example:** Artificial intelligence is the umbrella, machine learning learns patterns from data, deep learning uses multilayer neural networks and generative AI produces content; capability claims must identify the actual method and task. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication.

Qualification: The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Prediction-causation boundary. **Named evidence/example:** AI systems often generate probabilistic predictions or outputs from learned patterns; confidence and correlation do not establish verified truth, causation, fairness or legal justification. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Foundation-model boundary. **Named evidence/example:** A foundation model is pretrained for broad downstream adaptation, so data, licence, bias and safety failures can propagate to many applications; a model is not identical with every product built on it. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Hallucination-boundary. **Named evidence/example:** Generative AI can produce fluent false content because it predicts plausible continuations rather than retrieving verified truth; grounding, citations, human review and outcome monitoring remain necessary. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Executable exam-length answer / compression plan: For 10 marks, spend one-sixth of the time decoding the directive and drawing definition -> mechanism -> system/institution -> application/status -> risk/outcome; write four to seven claim -> named evidence -> analysis -> qualification points; reserve the final minute for unit, date, actor, platform, approval/test/deployment rung and residual-risk checks.

Why this earns marks: The answer obeys the directive, explains science rather than listing schemes, uses India-centric evidence and preserves technical, institutional, regulatory and maturity distinctions.

How to improve this answer: For "Distinguish AI, machine learning, deep learning and generative AI. Answer in about 150 words.", replace the weakest catalogue point with one exact mechanism, named mission/institution/platform, dated status, measurable boundary, implementation constraint and answer-specific qualification.

ORIGINAL MAINS 2 - 10 MARKS

Question: Explain why model capability is not the same as accountable deployment. Answer in about 150 words.

Model thesis: Claim: Capability-deployment boundary. **Named evidence/example:** A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Use-case-model boundary. **Named evidence/example:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Decision-support boundary. **Named evidence/example:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Claim -> named evidence -> analysis -> qualification:

- A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting.
- An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress.
- AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal.

Qualified conclusion: Claim: Capability-deployment boundary. **Named evidence/example:** A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication.

Qualification: The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Use-case-model boundary. **Named evidence/example:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication.

Qualification: The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Decision-support boundary. **Named evidence/example:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication.

Qualification: The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Demand decoding: The directive **explain** requires a direct position on "Explain why model capability is not the same as accountable deployment. Answer in about 150...", each clause, the exact technical mechanism, named Indian institution, dated status, application, risk and qualification.

Detailed examiner-grade model answer:

Introduction and thesis: Claim: Capability-deployment boundary. **Named evidence/example:** A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication.

Qualification: The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Use-case-model boundary. **Named evidence/example:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication.

Qualification: The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Decision-support boundary. **Named evidence/example:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication.

Qualification: The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Analytical body:

1. **Claim and named evidence:** A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.

2. **Claim and named evidence:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
3. **Claim and named evidence:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.

Counter-position / limit: A launch, test, target, budget, capacity, approval, contract, dataset, benchmark or prototype cannot alone establish deployment, induction, commercial operation, safety, inclusion or observed outcome; test mechanism, status, scale and evidence.

Qualified conclusion: Claim: Capability-deployment boundary. **Named evidence/example:** A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Use-case-model boundary. **Named evidence/example:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Decision-support boundary. **Named evidence/example:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Executable exam-length answer / compression plan: For 10 marks, spend one-sixth of the time decoding the directive and drawing definition -> mechanism -> system/institution -> application/status -> risk/outcome; write four to seven claim -> named evidence -> analysis -> qualification points; reserve the final minute for unit, date, actor, platform, approval/test/deployment rung and residual-risk checks.

Why this earns marks: The answer obeys the directive, explains science rather than listing schemes, uses India-centric evidence and preserves technical, institutional, regulatory and maturity distinctions.

How to improve this answer: For "Explain why model capability is not the same as accountable deployment. Answer in about 150...", replace the weakest catalogue point with one exact mechanism, named mission/institution/platform, dated status, measurable boundary, implementation constraint and answer-specific qualification.

ORIGINAL MAINS 3 - 15 MARKS

Question: Describe IndiaAI as a mission ecosystem while preserving pillar and implementation status. Answer in about 250 words.

Model thesis: Claim: IndiaAI-mission boundary. **Named evidence/example:** IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Implementation-institution boundary. **Named evidence/example:** MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion

retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Pillar-status boundary. **Named evidence/example:** A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Compute-measure boundary. **Named evidence/example:** A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** AIKosh-data boundary. **Named evidence/example:** AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Claim -> named evidence -> analysis -> qualification:

- IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code.
- MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making.
- A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence.
- A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition.
- AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access.

Qualified conclusion: **Claim:** IndiaAI-mission boundary. **Named evidence/example:** IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Implementation-institution boundary. **Named evidence/example:** MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Pillar-status boundary. **Named evidence/example:** A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Compute-measure boundary. **Named evidence/example:** A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** AIKosh-data boundary. **Named evidence/example:** AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Demand decoding: The directive **describe** requires a direct position on “Describe IndiaAI as a mission ecosystem while preserving pillar and implementation status....”, each clause, the exact technical mechanism, named Indian institution, dated status, application, risk and qualification.

Detailed examiner-grade model answer:

Introduction and thesis: **Claim:** IndiaAI-mission boundary. **Named evidence/example:** IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Implementation-institution boundary. **Named evidence/example:** MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Pillar-status boundary. **Named evidence/example:** A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Compute-measure boundary. **Named evidence/example:** A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** AIKosh-data boundary. **Named evidence/example:** AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Analytical body:

1. **Claim and named evidence:** IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
2. **Claim and named evidence:** MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
3. **Claim and named evidence:** A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
4. **Claim and named evidence:** A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.

5. **Claim and named evidence:** AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.

Counter-position / limit: A launch, test, target, budget, capacity, approval, contract, dataset, benchmark or prototype cannot alone establish deployment, induction, commercial operation, safety, inclusion or observed outcome; test mechanism, status, scale and evidence.

Qualified conclusion: Claim: IndiaAI-mission boundary. **Named evidence/example:** IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Implementation-institution boundary. **Named evidence/example:** MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Pillar-status boundary. **Named evidence/example:** A mission pillar identifies an implementation domain, not a completed output; approval, expression of interest, empanelment, procurement, model proposal, deployment and measured outcome require separate evidence. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Compute-measure boundary. **Named evidence/example:** A GPU objective, empanelled capacity, affordable compute unit, booked service and physically deployed accelerator are different measures; no number can be converted into another without an official definition. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** AIKosh-data boundary. **Named evidence/example:** AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Executable exam-length answer / compression plan: For 15 marks, spend one-sixth of the time decoding the directive and drawing definition -> mechanism -> system/institution -> application/status -> risk/outcome; write four to seven claim -> named evidence -> analysis -> qualification points; reserve the final minute for unit, date, actor, platform, approval/test/deployment rung and residual-risk checks.

Why this earns marks: The answer obeys the directive, explains science rather than listing schemes, uses India-centric evidence and preserves technical, institutional, regulatory and maturity distinctions.

How to improve this answer: For "Describe IndiaAI as a mission ecosystem while preserving pillar and implementation status...", replace the weakest catalogue point with one exact mechanism, named mission/institution/platform, dated status, measurable boundary, implementation constraint and answer-specific qualification.

ORIGINAL MAINS 4 - 15 MARKS

Question: Analyse lifecycle risks in public-sector AI and the controls needed at each stage. Answer in about 250 words.

Model thesis: Claim: Use-case-model boundary. **Named evidence/example:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Claim: Decision-support boundary. **Named evidence/example:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Claim: Risk-lifecycle boundary. **Named evidence/example:** AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:**

Bias-fairness boundary. **Named evidence/example:** Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Claim: Privacy-security boundary. **Named evidence/example:** Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Claim -> named evidence -> analysis -> qualification:

- An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress.
- AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal.
- AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle.
- Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy.
- Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability.

Qualified conclusion: Claim: Use-case-model boundary. **Named evidence/example:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Decision-support boundary. **Named evidence/example:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Risk-lifecycle boundary. **Named evidence/example:** AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage

cannot control the full lifecycle. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Bias-fairness boundary. **Named evidence/example:** Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Privacy-security boundary. **Named evidence/example:** Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Demand decoding: The directive **analyse** requires a direct position on "Analyse lifecycle risks in public-sector AI and the controls needed at each stage. Answer in...", each clause, the exact technical mechanism, named Indian institution, dated status, application, risk and qualification.

Detailed examiner-grade model answer:

Introduction and thesis: **Claim:** Use-case-model boundary. **Named evidence/example:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Decision-support boundary. **Named evidence/example:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Risk-lifecycle boundary. **Named evidence/example:** AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Bias-fairness boundary. **Named evidence/example:** Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Privacy-security boundary. **Named evidence/example:** Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Analytical body:

1. **Claim and named evidence:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
2. **Claim and named evidence:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway ->

public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.

3. **Claim and named evidence:** AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
4. **Claim and named evidence:** Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
5. **Claim and named evidence:** Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.

Counter-position / limit: A launch, test, target, budget, capacity, approval, contract, dataset, benchmark or prototype cannot alone establish deployment, induction, commercial operation, safety, inclusion or observed outcome; test mechanism, status, scale and evidence.

Qualified conclusion: **Claim:** Use-case-model boundary. **Named evidence/example:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Decision-support boundary. **Named evidence/example:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Risk-lifecycle boundary. **Named evidence/example:** AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Bias-fairness boundary. **Named evidence/example:** Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Privacy-security boundary. **Named evidence/example:** Privacy governs lawful data use and purpose, while security protects systems and data from compromise; both matter, but neither alone resolves bias, opacity, safety, misinformation or accountability. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Executable exam-length answer / compression plan: For 15 marks, spend one-sixth of the time decoding the directive and drawing definition -> mechanism -> system/institution -> application/status -> risk/outcome; write four to seven claim -> named evidence -> analysis -> qualification points; reserve the final minute for unit, date, actor, platform, approval/test/deployment rung and residual-risk checks.

Why this earns marks: The answer obeys the directive, explains science rather than listing schemes, uses India-centric evidence and preserves technical, institutional, regulatory and maturity distinctions.

How to improve this answer: For “Analyse lifecycle risks in public-sector AI and the controls needed at each stage. Answer in...”, replace the weakest catalogue point with one exact mechanism, named mission/institution/platform, dated status, measurable boundary, implementation constraint and answer-specific qualification.

ORIGINAL MAINS 5 - 20 MARKS

Question: Evaluate India's AI-governance approach through legal instruments, sectoral duties and mission capability. Answer in about 300 words.

Model thesis: Claim: IndiaAI-mission boundary. **Named evidence/example:** IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Implementation-institution boundary. **Named evidence/example:** MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Deepfake-instrument boundary. **Named evidence/example:** Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Guideline-law boundary. **Named evidence/example:** An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Sectoral-governance boundary. **Named evidence/example:** AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Volatile-mission boundary. **Named evidence/example:** Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Claim -> named evidence -> analysis -> qualification:

- IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code.
- MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making.
- Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime.
- An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record.
- AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators.

- Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability.

Qualified conclusion: Claim: IndiaAI-mission boundary. **Named evidence/example:** IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Implementation-institution boundary. **Named evidence/example:** MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Deepfake-instrument boundary. **Named evidence/example:** Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Guideline-law boundary. **Named evidence/example:** An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Sectoral-governance boundary. **Named evidence/example:** AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Volatile-mission boundary. **Named evidence/example:** Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Demand decoding: The directive **evaluate** requires a direct position on "Evaluate India's AI-governance approach through legal instruments, sectoral duties and...", each clause, the exact technical mechanism, named Indian institution, dated status, application, risk and qualification.

Detailed examiner-grade model answer:

Introduction and thesis: Claim: IndiaAI-mission boundary. **Named evidence/example:** IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Implementation-institution boundary. **Named evidence/example:** MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Deepfake-instrument boundary. **Named evidence/example:** Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment

status boundary. **Claim:** Guideline-law boundary. **Named evidence/example:** An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Sectoral-governance boundary. **Named evidence/example:** AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Volatile-mission boundary. **Named evidence/example:** Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Analytical body:

1. **Claim and named evidence:** IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
2. **Claim and named evidence:** MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
3. **Claim and named evidence:** Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
4. **Claim and named evidence:** An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
5. **Claim and named evidence:** AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
6. **Claim and named evidence:** Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.

Counter-position / limit: A launch, test, target, budget, capacity, approval, contract, dataset, benchmark or prototype cannot alone establish deployment, induction, commercial operation, safety, inclusion or observed outcome; test mechanism, status, scale and evidence.

Qualified conclusion: Claim: IndiaAI-mission boundary. **Named evidence/example:** IndiaAI is an ecosystem mission spanning compute, innovation, applications, AIKosh, startup finance, FutureSkills and Safe and Trusted AI; it is not a single model, company, subsidy or legal code. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Implementation-institution boundary. **Named evidence/example:** MeitY anchors policy and the IndiaAI Mission is implemented through IndiaAI within Digital India Corporation; NITI Aayog's earlier strategy role is distinct from mission implementation and rule-making. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Deepfake-instrument boundary. **Named evidence/example:** Synthetic-media harms may be addressed through intermediary rules, labelling, provenance and enforcement, but a rule on distribution is not the same as a comprehensive AI statute or model-safety regime. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Guideline-law boundary. **Named evidence/example:** An advisory, consultation report or governance guideline is non-binding unless supported by enforceable law, while notified rules under a parent Act create subordinate legislation; India has no standalone AI Act in the owner record. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Sectoral-governance boundary. **Named evidence/example:** AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Volatile-mission boundary. **Named evidence/example:** Mission outlay, compute count, model selection, model status, safety-institute status and summit outcomes require dated official IndiaAI or MeitY sources; announcements cannot be rewritten as deployed capability. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Executable exam-length answer / compression plan: For 20 marks, spend one-sixth of the time decoding the directive and drawing definition -> mechanism -> system/institution -> application/status -> risk/outcome; write four to seven claim -> named evidence -> analysis -> qualification points; reserve the final minute for unit, date, actor, platform, approval/test/deployment rung and residual-risk checks.

Why this earns marks: The answer obeys the directive, explains science rather than listing schemes, uses India-centric evidence and preserves technical, institutional, regulatory and maturity distinctions.

How to improve this answer: For "Evaluate India's AI-governance approach through legal instruments, sectoral duties and...", replace the weakest catalogue point with one exact mechanism, named mission/institution/platform, dated status, measurable boundary, implementation constraint and answer-specific qualification.

ORIGINAL MAINS 6 - 20 MARKS

Question: Design an accountable public procurement framework for high-impact AI systems. Answer in about 300 words.

Model thesis: Claim: Capability-deployment boundary. **Named evidence/example:** A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** AIKosh-data boundary. **Named evidence/example:** AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. **Analysis:** This fixes the scientific process, system boundary, institution,

application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Claim: Use-case-model boundary. **Named evidence/example:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Decision-support boundary. **Named evidence/example:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Risk-lifecycle boundary. **Named evidence/example:** AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Bias-fairness boundary. **Named evidence/example:** Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Procurement-accountability boundary. **Named evidence/example:** Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Sectoral-governance boundary. **Named evidence/example:** AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Claim -> named evidence -> analysis -> qualification:

- A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting.
- AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access.
- An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress.
- AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal.
- AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle.
- Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy.
- Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability.
- AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators.

Qualified conclusion: Claim: Capability-deployment boundary. **Named evidence/example:** A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** AIKosh-data boundary. **Named evidence/example:** AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Use-case-model boundary. **Named evidence/example:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Decision-support boundary. **Named evidence/example:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Risk-lifecycle boundary. **Named evidence/example:** AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Bias-fairness boundary. **Named evidence/example:** Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Procurement-accountability boundary. **Named evidence/example:** Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Sectoral-governance boundary. **Named evidence/example:** AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Demand decoding: The directive **answer** requires a direct position on "Design an accountable public procurement framework for high-impact AI systems. Answer in...", each clause, the exact technical mechanism, named Indian institution, dated status, application, risk and qualification.

Detailed examiner-grade model answer:

Introduction and thesis: Claim: Capability-deployment boundary. **Named evidence/example:** A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** AIKosh-data boundary. **Named evidence/example:** AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset

representativeness, lawful provenance, suitability, quality or complete access. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Use-case-model boundary. **Named evidence/example:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Decision-support boundary. **Named evidence/example:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Risk-lifecycle boundary. **Named evidence/example:** AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Bias-fairness boundary. **Named evidence/example:** Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Procurement-accountability boundary. **Named evidence/example:** Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Sectoral-governance boundary. **Named evidence/example:** AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Analytical body:

1. **Claim and named evidence:** A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
2. **Claim and named evidence:** AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
3. **Claim and named evidence:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
4. **Claim and named evidence:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and

appeal. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.

5. **Claim and named evidence:** AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
6. **Claim and named evidence:** Bias may arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
7. **Claim and named evidence:** Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.
8. **Claim and named evidence:** AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. **Analysis:** Connect mechanism -> named system/institution -> application or implementation pathway -> public or strategic consequence. **Qualification:** State platform, unit, source/date/status, maturity, regulatory limit, uncertainty, exception or residual risk.

Counter-position / limit: A launch, test, target, budget, capacity, approval, contract, dataset, benchmark or prototype cannot alone establish deployment, induction, commercial operation, safety, inclusion or observed outcome; test mechanism, status, scale and evidence.

Qualified conclusion: Claim: Capability-deployment boundary. **Named evidence/example:** A model that demonstrates a capability or performs well on a benchmark is not automatically a validated, integrated or accountable deployment in health, welfare, policing, credit or another high-impact setting. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** AIKosh-data boundary. **Named evidence/example:** AIKosh is the mission's dataset and model-resource platform layer; platform availability does not prove dataset representativeness, lawful provenance, suitability, quality or complete access. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Use-case-model boundary. **Named evidence/example:** An AI use case describes a problem and workflow, while a model is one component and deployment includes data pipelines, interfaces, human decisions, security, monitoring and grievance redress. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Decision-support boundary. **Named evidence/example:** AI can assist a public official or clinician without lawfully replacing accountable judgement; automated adverse decisions affecting rights require reasons, human oversight, contestability and appeal. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Risk-lifecycle boundary. **Named evidence/example:** AI risk enters through problem selection, data, training, evaluation, procurement, integration, use and post-deployment change; a generic ethics statement at the final stage cannot control the full lifecycle. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Bias-fairness boundary. **Named evidence/example:** Bias may

arise from representation, labels, proxies, optimisation or deployment context, and equal aggregate accuracy can conceal unequal group harm; fairness requires context-specific testing and remedy. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Procurement-accountability boundary. **Named evidence/example:** Public procurement must specify data rights, evaluation, audit logs, human oversight, incident reporting, vendor change control and exit obligations; buying an AI service does not transfer constitutional accountability. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary. **Claim:** Sectoral-governance boundary. **Named evidence/example:** AI risk differs across health, finance, education, policing, welfare and low-risk productivity tools, so proportionate governance should combine horizontal principles with sectoral law and regulators. **Analysis:** This fixes the scientific process, system boundary, institution, application and verified capability rung before drawing the policy implication. **Qualification:** The conclusion retains the owner's date, unit, institution, system category, legal character and development-to-deployment status boundary.

Executable exam-length answer / compression plan: For 20 marks, spend one-sixth of the time decoding the directive and drawing definition -> mechanism -> system/institution -> application/status -> risk/outcome; write four to seven claim -> named evidence -> analysis -> qualification points; reserve the final minute for unit, date, actor, platform, approval/test/deployment rung and residual-risk checks.

Why this earns marks: The answer obeys the directive, explains science rather than listing schemes, uses India-centric evidence and preserves technical, institutional, regulatory and maturity distinctions.

How to improve this answer: For "Design an accountable public procurement framework for high-impact AI systems. Answer in...", replace the weakest catalogue point with one exact mechanism, named mission/institution/platform, dated status, measurable boundary, implementation constraint and answer-specific qualification.